WM-4x recognition training (Qwen3-1.7B, 3 epochs)

768 watermarked requests; whole model + classifier head. Logged once per epoch (96 steps each).

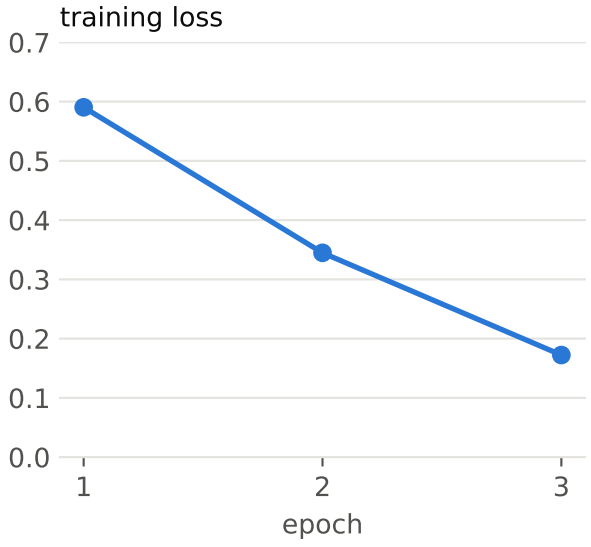

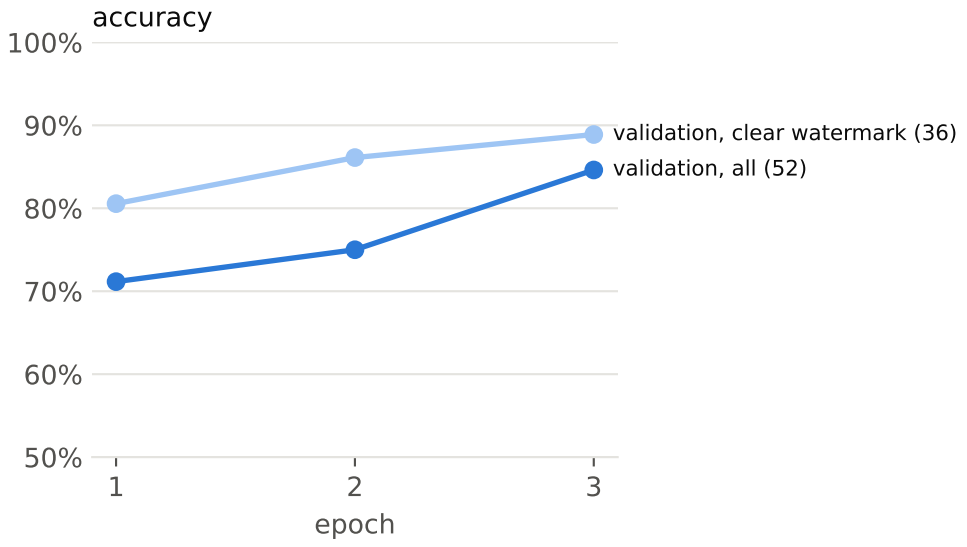

Epoch 3 kept (best validation). Held-out test: 90.4% (52 texts), 100% on the 28 clearly watermarked ones. Small evaluation sets.